

# CS 61 Lecture 12

Thursday, Oct 10<sup>th</sup>, 2009

## Kernel Lecture #2

### Allowed on midterm:

- internet to access website, previous yrs of the website
- open book, open notes
- assembly manual

### Patch

- a difference file

~~@@ -11,6 +11,11 @@ void process\_main()~~  
 added 5 lines starting at line 6

~~@@ -137,8 +137,8 @@ void process\_setup(...)~~

✓ replaced some # of lines

a) kernel / kernel.cc  
 b) kernel / kernel.c

- applying a diff file w/ patch

patch < patch09-kernel-knum.diff

### Last Time

- initinil loop executed in the context of kernel by modifying its code

virt8-t\* syscall = (virt8-t\*) 0xheadb;  
 fabricate a pointer to memory

syscalls[0] = 0xdad; ? while (true)

$\text{SYSCALL}[1] = 0x\text{fc}$ ] infinite loop

`lwnut->printf(0x0E0D, "%d", sys_getpid());`

infinite loop in x86-64

-b  $\text{fc}$

→ signed 8 bit int = 0b 1111 1110 = -2

$-X = \sim X + 1$

flipping all bits +1

negative of  $\text{fc} = 0b000\ 0010 = 2$

program

executes next line, -2 = current line

Where does the LL go?

`0x400ad6` → corresponds to `syscall_entry` function

What if EIP jumped to this syscall address

further than writing and invoking?

`void (*f)() = (void(*)()) syscalls;`

`f();`

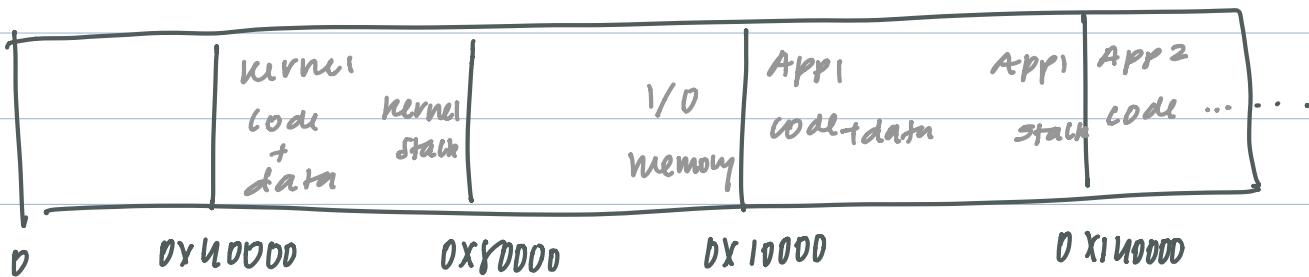
\* code -eve has compiled to: less obj/p-eve.asm

`ff d0 callg *rax`

↳ executing kernel code := changing privilege  
of which eve is running

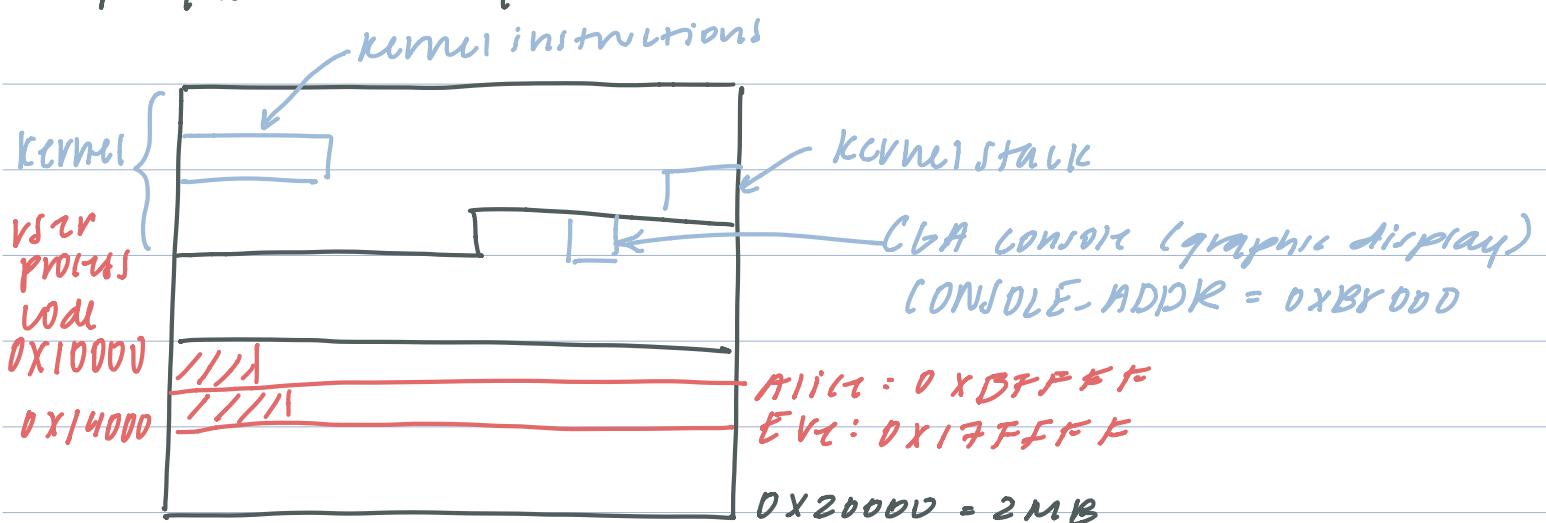
eve gains privilege when kernel gives it

# X86-64 Memory Layout



# Weniny DS Memory Layout

## Physical memory



```

for (vmiter it (kernel_pagetable, 0));
    iterate over { it.va() < MEMSIZE_PHYSICAL;
        pages (212 bytes) } it += PAGE_SIZE {
            if (it.va() == 0) { maps to itself
                allow anyone to access it
                → it.map(it.va(), PTE_P | PTE_W | PTE_U);
                cannot dereference null pointer
                it.map(it.va(), 0); }
            }
    }
}

```

initializes kernel memory mappings for physical memory

{ add: if(it.val() >= PROL-START-ADDR) {

    it.map(it.val(), PTE-P|PTE-W, PTEU);

}

remove PTE-U from previous

→ still causes exception in (page fault)

we is trying to access buffer (on console  
page)

change to:

if(it.val() >= PROL-START-ADDR ||

    it.val() == CONSOLE-ADDR)

allow write/mem  
access to console

    it.map(it.val(), PTE-P|PTE-W, PTEU);

}

Protected control transfer → SLOW.

processors mechanisms for transferring control

across privilege domains (less → more privileged)

X86-64

- Kernel configures entry points for p.c.t.
- one entry point for every interrupt, fault,  
and trap
- one entry point for syscall instruction

## interrupt / fault / trap

- kernel configures stack for protected control transfer at boot time
- when interrupt / fault / trap occurs:
  1. processor changes to configured stack
  2. processor pushes old stack pointer onto new stack + other critical regs (i·rip,  
has privilege bits → i·cs)  
*value of instruction that caused fault*
  3. processor installs the kernel's configured entrypoint and i·cs register ← changes privilege at entrypoint  
*changes instruction that is excluded next*
  4. kernel takes control  
(K exception.s)

irretz : undos work (return to process)

↳ values loaded from process descriptor

## skip over faulting instruction

regs → reg-rip += 8  
→ faulting instruction is 8 bytes long

## kernel's configured entrypoint

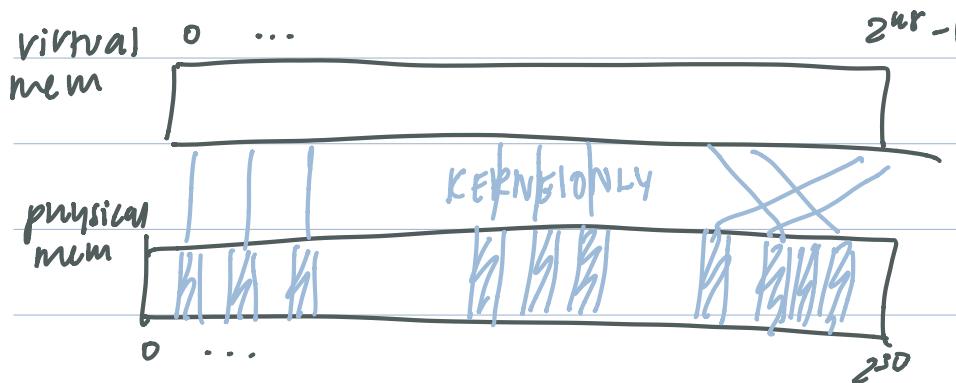
jump to syscalls - mny upon initialization

## Kill the process

currnt → stat = P-BROKEN

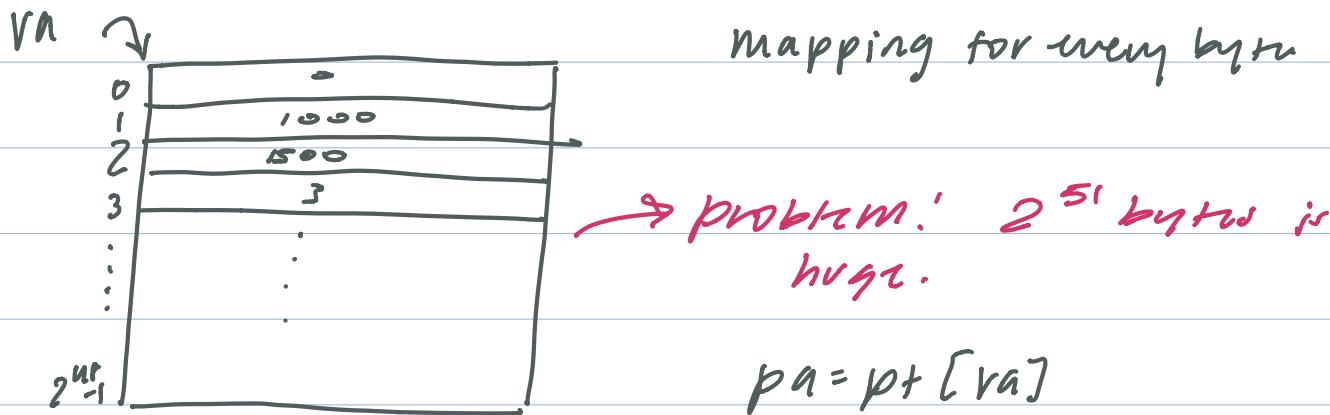
# Implementation of virtual memory

## X86-64 page tables



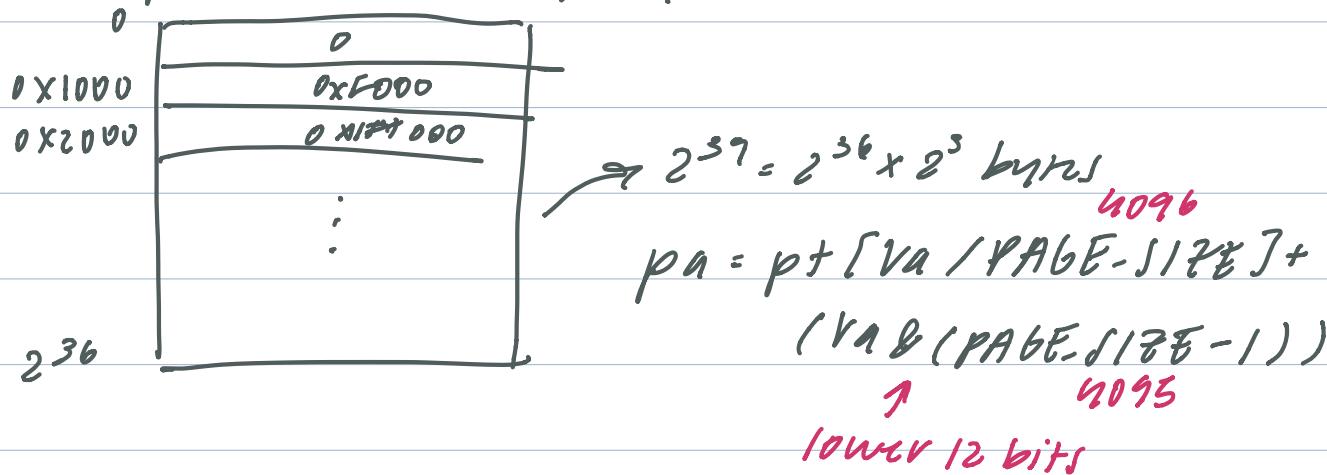
KERNONLY sets up  
processor implements

## Strawman: Byte map



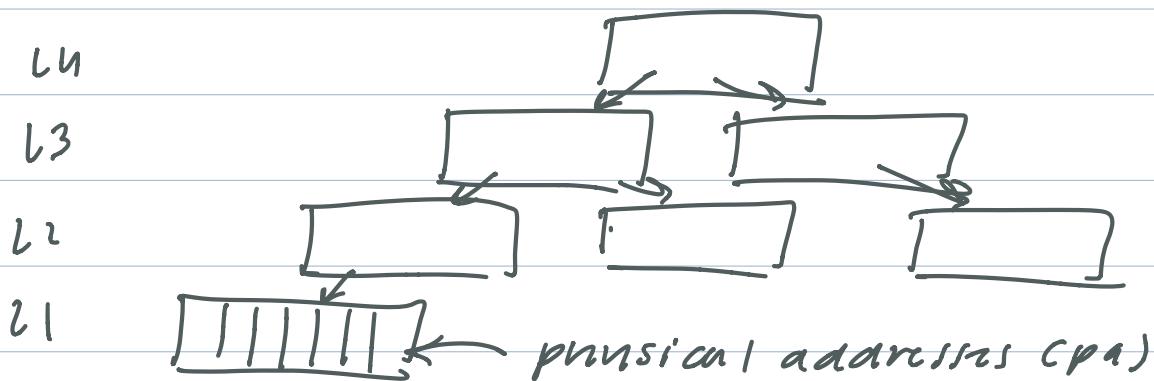
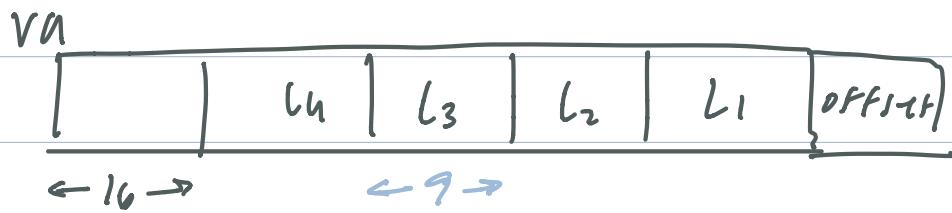
## Strawman: Page Map

Map blocks ( $2^{12}$  bytes), aka "pages"



## Try design (compress out 0s)

Compress address space: process will use  $2^{16}$  bytes



$2^{14}$  bytes in W-Entry (instead of  $2^{39}$ )