

# CS61 LECTURE 16

Thursday, Oct 24

## New Unit: Storage

memory, disk

SSD, hard drive

iCloud

cache

moving from

very simple model of memory

address  $\leftrightarrow$  read/write



memory system

hierarchy of storage devices

w/ different capacities, costs,

and times to access

### Analogy: Books

intrad of data

|               |            |                           |        |
|---------------|------------|---------------------------|--------|
| disk          | 3 books    | = registers               | ~16    |
| shelves       | 60 books   | = cache                   |        |
| Widener       | 3.5M books | = main memory             | 32GB   |
| Depository    | 10M books  | = disk drive              | 256 GB |
| Borrow Direct | 90M books  | = network / cloud storage |        |

heirarchy

- Why not connect the processor directly to main memory?
- add A, B not possible, why? problems w/ costs and time to access*

appli. com

volatil. erased when power off

DDR4 DRAM

16 GB \$400 \$25/GB 12.5-15 ns

SSD/flash

256 GB \$200 \$0.80/GB 0.1 ms

Hard Drive

2TB \$100 \$0.50/GB 12 ms

non-volatile, still present when power off

## KEY OBSERVATION:

CAN'T BUILD STORAGE THAT IS BIG, FAST, AND CHEAP  $\rightarrow$  create a system w/ appearance of big and fast

## HOW? books

- ① most likely to read
- ② things I generate (wrote)

data

## Principle: Locality of reference

- good (fast on modern machines) take advantage of locality of reference).

### 1. Temporal locality

if referenced, likely to reference again soon

### 2. Spatial locality

if referenced, those around it likely referenced

```
long f(const int* V, int n) {
    long sum = 0;
    for (int i = 0; i < n; ++i) {
        sum += V[i];
    }
}
```

temporal locality

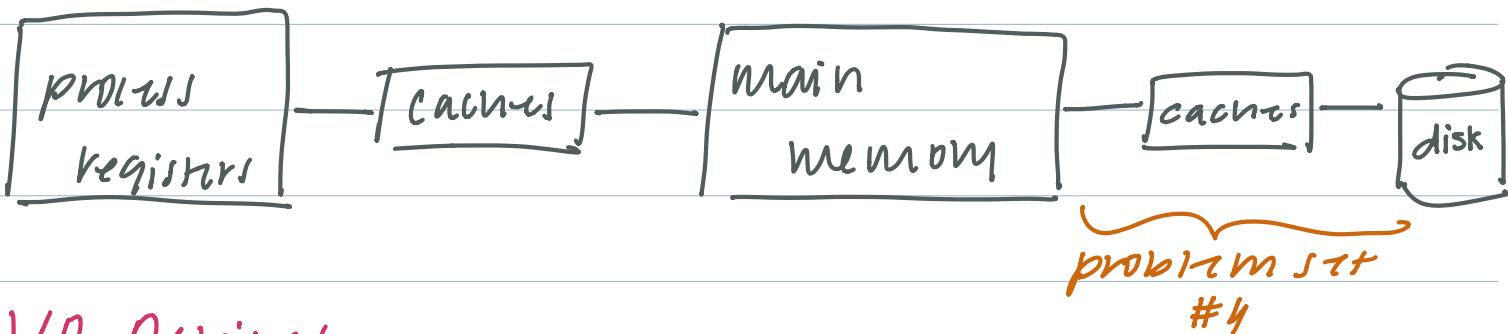
spatial locality

\* instructions themselves have temporal and spatial locality

## Key: Caching (verb)

using small amount of fast storage to speed up access  
to some large, slow storage

simple idea, many variants → popular careers



## I/O Devices

VMS file = sequence of bytes

↳ treats every I/O device like a file

\* some OS have different end-line characters, example <sup>for</sup>

- files as an abstraction of I/O devices
- virtual memory as an abstraction of main memory
- processes as an abstraction for running programs

```
int main()
```

```
int fd = STDOUT_FILENO;
```

```
if (isatty(fd)) {
```

```
    fd = open(DATAFILE, O_WRONLY | O_CREAT);
```

```
}
```

```
if (fd < 0) {
```

```
    perror("open");
```

```
    exit(1);
```

filename

with  
only

create if  
doesn't  
exist

```
O_TRUNC | O_SYNC, 0600);
```

if file exist,  
throw it out

don't allow  
prog to cont.

permissions  
for everyone  
to read/write

until program  
finishes writing

{

5.12 MB

ssize\_t size = 5120000;

const char\* buf = "G";

double start = tstamp();

while (n &lt; size) {

ssize\_t r = write(fd, buf, 1);

if (r != 1){

perror("write");

exit(1);

{

n += r;

if (h-&gt;PRINT\_FREQUENCY == 0){

report(n, tstamp() - start);

{

{

close(fd);

report(n, tstamp() - start);

fprintf(stderr, "\n");

{

Original

 $\approx 55 \text{ kB of memory / second}$   
 $= 0.4 \text{ ms / byte}$ 

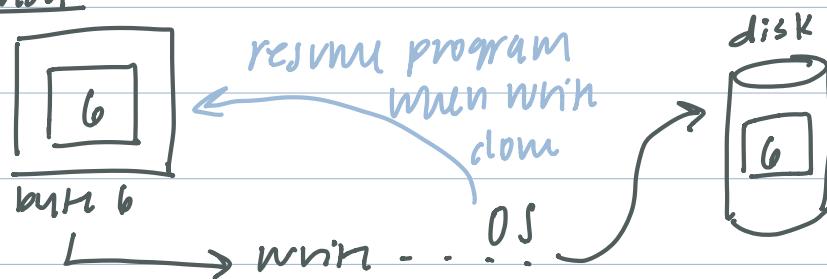
remove O\_SYNC flag

 $\approx 9.2 \text{ kB / second}$ use stdio (standard i/o library)  $\leftarrow$  fwrite $\approx 63.2 \text{ kB / second}$ 

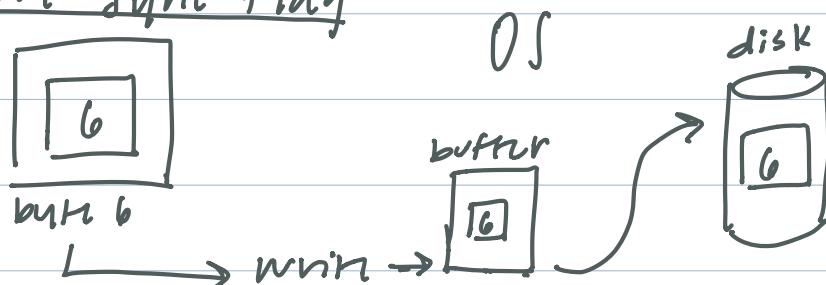
fwrite

# program WDX-Y

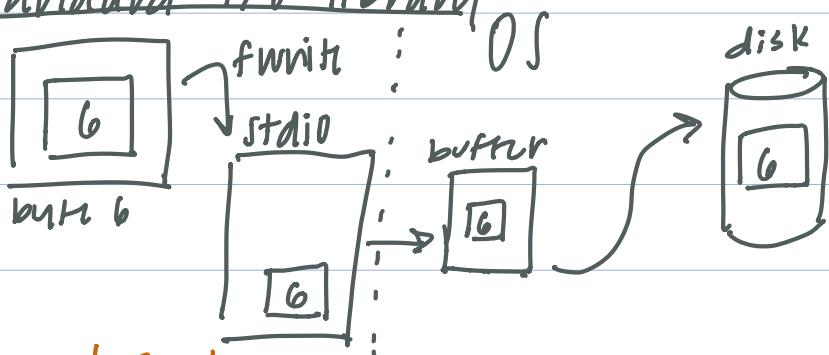
## original



## remove sync flag



## standard i/o library



essentially doing a store