


CS61 Lecture 18: Storage 3

October 31,

Continue designing & modeling caches 2019

single slot \rightarrow multislot caches

Multislot Cache



params

- int nslots (number of slots in cache)
- bool is-direct-mapped
- char replacement (replacement algo if fully associative)
- cache_slot[]
 - valid (false if empty)
 - mindx (index into data array, similar to address)
 - last-access (helps w/ replacement)


Replacement Scheme

1. Random
2. Least Recently used

Direct-Mapped Cache

ref A \rightarrow f

only one slot that
a certain reference
can go into




Fully-Associative Cache

ref A \rightarrow f

could be put
anywhere

Summary: How caches \uparrow perform

1. PREFETCHING



bring data closer before it's needed

- eg: prefetch file from disk when opened
works well when user ops are anticipated correctly

2. BATCHING



change a large number of small requests
into small number of large requests

- eg: sequentially reading 4096 bytes from disk
- eg: stdio cache and file cache faster by writing in chunks

works well when per-request cost R is high
and per-unit cost U is low

$$C = NU + R$$

locality benefits

3. WRITE COALESCING



eliminate all but last write

• eg: index variable in loop

works well for technologies where write is
complex (ex: flash memory)

Unix File (a sequence of bytes)

1. ones that you can read like a book

RANDOM ACCESS FILES in Unix

2. ones that are like listening to your grandmother
tell a story

STREAMS in Unix

RAM File

length	finite
skip around	yes (file seekable)
capture in memory?	yes (mappable)

seek call

map call

Stream File

possibly infinite
no (not seekable)

no (not mappable)

File Descriptor

includes / remembers the current file position

read (fd, buf, len)

write (fd, buf, len)

lseek (