


strace

- debugs syscalls by another program

`strace -o strace.out PROGRAMNAME ARGS ...`


Access patterns

Sequential access




cache 4 datapoints at the end

Reverse-sequential access



align to 4-bit groups

Strided Access



- access data w/ uniform skip
- need large cache

- if user tries to read > size of cache, have to make multiple reads and/or use storage

strat exercise

`lseek(0, 4114, SEEK_SET);` how to jump to new address


strat 05

`read(0, "xxxx "..., 4096) = 4091`
↓ ↗
requested # of bytes to read actually read # of bytes

used cache

Single-slot cache

`struct io61_file {` ← \equiv to `FILE*` from `fopen()` in standard i/o library
 `int fd;` ← file descriptor
 `static constexpr off_t bufsize = 4096;` // or whatever
 `unsigned char cbuf[bufsize];` ↗ can change
 `off_t tag;` // file offset of first byte in cache (0 when file is opened)
 `off_t end_tag;` // file offset one past last valid byte in cache
 `off_t pos_tag;` // file offset of next char to read in cache
};



checks

* assuming one byte

• is the byte inside the buffer or not

keep track of byte # of beginning and end

Facts about cache representation

- reads or writes but not both
- $tag \leq end_tag$
- $end_tag - tag \leq \text{bufsize}$
not always equal bc end/start tag can move
- if $tag == end_tag$

cache slot is empty (no valid data)

- $tag \leq end_tag$ and $pos_tag \leq end_tag$
read write

current file pos = end_tag tag

- the next `io61-read` or `io61-write` call starts at offset pos_tag .
- if $i \geq tag$ and $i \leq end_tag$, `cbuf[i-tag] = byte at offset i`

Filling cache

```
void io61_fill(io61_file* f) {  
    // Fill the read cache with new data, starting from file offset `end_tag`.  
    // Only called for read caches.  
  
    // Check invariants.  
    assert(f->tag <= f->pos_tag && f->pos_tag <= f->end_tag);  
    assert(f->end_tag - f->pos_tag <= f->bufsize);  
  
    /* ANSWER */  
    // Reset the cache to empty.  
    f->tag = f->pos_tag = f->end_tag;  
    // Read data. read (file descriptor, buffer, buffer size)  
    ssize_t n = read(f->fd, f->cbuf, f->bufsize);  
    if (n >= 0) { unix read  
        f->end_tag = f->tag + n;  
    }  
    n is -1 if failed read  
    // Recheck invariants (good practice!).  
    assert(f->tag <= f->pos_tag && f->pos_tag <= f->end_tag);  
    assert(f->end_tag - f->pos_tag <= f->bufsize);  
}
```

Easy read

```
ssize_t io61_read(io61_file* f, char* buf, size_t sz) {
    // Check invariants.
    assert(f->tag <= f->pos_tag && f->pos_tag <= f->end_tag);
    assert(f->end_tag - f->pos_tag <= f->bufsize);

    // assuming hit.
    // The desired data is guaranteed to lie within this cache slot.
    assert(sz <= f->bufsize && f->pos_tag + sz <= f->end_tag);

    /* ANSWER */
    ① memcpy(buf, &f->cbuf[f->pos_tag - f->tag], sz);
    ② f->pos_tag += sz;
}
```

1. calculate index in cbuf and read that index
2. update pos_tag

Full Read

```
ssize_t io61_read(io61_file* f, char* buf, size_t sz) {
    // Check invariants.
    assert(f->tag <= f->pos_tag && f->pos_tag <= f->end_tag);
    assert(f->end_tag - f->pos_tag <= f->bufsize);

    /* ANSWER */
    size_t pos = 0;
    while (pos < sz) {
        if (f->pos_tag == f->end_tag) {
            io61_fill(f);
            if (f->pos_tag == f->end_tag) {
                break; or return pos;
            }
        }

        // This would be faster if you used `memcpy`!
        buf[pos] = f->cbuf[f->pos_tag - f->tag];
        ++f->pos_tag;
        ++pos;
    }
    return pos;
}
```

1. loop that reads a byte at a time cbuf -> buf

Easy write

```
ssize_t io61_write(io61_file* f, const char* buf, size_t sz) {  
    // Check invariants.  
    assert(f->tag <= f->pos_tag && f->pos_tag <= f->end_tag);  
    assert(f->end_tag - f->pos_tag <= f->bufsize);  
  
    // Write cache invariant.  
    assert(f->pos_tag == f->end_tag);  
  
    // The desired data is guaranteed to lie within this cache slot.  
    assert(sz <= f->bufsize && f->pos_tag + sz <= f->tag + f->bufsize);  
  
    /* ANSWER */  
    memcpy(&f->cbuf[f->pos_tag - f->tag], buf, sz);  
    f->pos_tag += sz;  
    f->end_tag += sz;  
    return sz;  
}
```